

Ray Qin

Senior Software & AI Engineer • B.S. Math-CS & Mechanical Engineering • Ex-ZipRecruiter (10M+ users)

San Francisco Bay Area • +1 415-604-6032 • qinjianxyz@gmail.com • linkedin.com/in/qinjianxyz • qinjianxyz.github.io

SUMMARY

Ships production AI end to end: autonomous agents, RAG, grounded LLM systems, and the full-stack products around them. At ZipRecruiter, rebuilt authentication for **10M+ monthly users**, lifting login conversion by double digits and cutting homepage load time about 40%, and led the team's move to AI-assisted delivery. Dual B.S. in Mathematics-Computer Science and Mechanical Engineering, UC San Diego.

EXPERIENCE

ZipRecruiter

Apr 2023 – Apr 2026

Senior Software Engineer • Santa Monica, CA • Remote

- Rebuilt registration, login, and authentication across **30+ fragmented entry points** into one reusable platform for 10M+ monthly users — a **double-digit login-conversion lift**, 10K+ weekly SMS opt-ins, and about **40% faster** homepage loads.
- Architected the backend (Go services, OTP login, auto-login tokens, Protobuf credential APIs) and shared React / Next.js frontend; sequenced the migration that retired three legacy auth, password-generation, and login services.
- Pioneered **AI-assisted delivery**: ran Claude Code loops that turned designs into reviewed PRs; wired Jira, Slack, Notion, and GitHub to agents via MCP; ran the company's first Cursor workshop and shaped the company's AI-usage guidelines.
- Shipped **vector-embedding similarity search** powering Suggested Jobs and More Candidates Like This, serving recommendations across the **10M+-user** marketplace.
- Drove cross-team delivery across messaging, mobile, SEO, paid acquisition, and identity; gave company-wide talks on API design and secure engineering.

INDEPENDENT WORK

Enterprise OS — multi-agent engineering system

Rust, agent orchestration, RAG

- Built a multi-agent system where Codex and Claude Code ship against real repos in isolated worktrees, every change gated by specs, tests, evals, scoped permissions, and CI; routinely lands **multi-file PRs autonomously**, reproducible and rollback-aware.
- Layered RAG over each codebase for grounded retrieval, with an eval harness that blocks regressions before merge and reduces human review to a final approval gate.

EARLIER EXPERIENCE

- **Eightfold** — SWE Intern (2022): made ETL loading atomic, cutting data inconsistency below 2%; extended consistency monitoring to 56 tables (RDS/Redshift/Firehose); built a React/Flask Spark console that cut debug time about 40%.
- **MEDsmart** — SWE Intern (2021–22): cut clinic operating cost about 11% and inventory waste about 4% with AWS data pipelines and Python/LSTM forecasting.
- **Qualcomm Institute (UCSD Calit2)** — SWE Intern (2020–21): built business-data visualizations from the Yelp API and city open data.

EDUCATION

University of California, San Diego

2019 – 2023

Dual B.S. — Mathematics-Computer Science & Mechanical Engineering • Cum Laude, GPA 3.9/4.0

SKILLS

Languages TypeScript / JavaScript, Go, Python, Rust, SQL, C / C++

AI LLM agents, RAG, retrieval & reranking, function calling, structured output, model routing, evals, MCP, vector search, context engineering; Claude Code, Codex, Cursor

Web & API React, Next.js, React Native, Node.js, gRPC / Connect / Protobuf, REST, auth / OTP, SEO

Infra & Data Postgres, Redis, Kafka, Docker, Kubernetes, Terraform, Spark, observability, TDD